

# Python Lab Exercise #2

## Objectives:

- Load .csv files into `pandas` DataFrames
- Describe and manipulate data in Series and DataFrames
- Visualize data using DataFrame methods and `matplotlib`

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## What is Pandas?

Pandas, as [the Anaconda docs](#) tell us, offers us "High-performance, easy-to-use data structures and data analysis tools." It's something like "Excel for Python", but it's quite a bit more powerful.

Let's read in the heart dataset.

Pandas has many methods for reading different types of files. Note that here we have a .csv file.

Read about this dataset [here](#).

```
heart_df = pd.read_csv('heart.csv')
```

The output of the `.read_csv()` function is a pandas *DataFrame*, which has a familiar tabular structure of rows and columns.

```
type(heart_df)
```

```
pandas.core.frame.DataFrame
```

```
heart_df
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
0	63	1	3	145	233	1	0	150	0
1	37	1	2	130	250	0	1	187	0
2	41	0	1	130	204	0	0	172	0
3	56	1	1	120	236	0	1	178	0
4	57	0	0	120	354	0	1	163	1
...	...	...	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1
299	45	1	3	110	264	0	1	132	0
300	68	1	0	144	193	1	1	141	0
301	57	1	0	130	131	0	1	115	1
302	57	0	1	130	236	0	0	174	0

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
...	...	...	...	...
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```
[303 rows x 14 columns]
```

## DataFrames and Series

Two main types of pandas objects are the DataFrame and the Series, the latter being in effect a single column of the former:

```
age_series = heart_df['age']
type(age_series)

pandas.core.series.Series
```

Notice how we can isolate a column of our DataFrame simply by using square brackets together with the name of the column.

Both Series and DataFrames have an *index* as well:

```
heart_df.index
RangeIndex(start=0, stop=303, step=1)

age_series.index
RangeIndex(start=0, stop=303, step=1)
```

Pandas is built on top of NumPy, and we can always access the NumPy array underlying a DataFrame using `.values`.

```
heart_df.values
array([[63., 1., 3., ..., 0., 1., 1.],
       [37., 1., 2., ..., 0., 2., 1.],
       [41., 0., 1., ..., 0., 2., 1.],
       ...,
       [68., 1., 0., ..., 2., 3., 0.],
       [57., 1., 0., ..., 1., 3., 0.],
       [57., 0., 1., ..., 1., 2., 0.]])
```

## Basic DataFrame Attributes and Methods

`.head()`

```
heart_df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
.tail()
```

```
heart_df.tail()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
oldpeak \									
298	57	0	0	140	241	0	1	123	1
0.2									
299	45	1	3	110	264	0	1	132	0
1.2									
300	68	1	0	144	193	1	1	141	0
3.4									
301	57	1	0	130	131	0	1	115	1
1.2									
302	57	0	1	130	236	0	0	174	0
0.0									

	slope	ca	thal	target
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```
.info()
```

```
heart_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   int64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
```

```
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
.describe()
```

```
heart_df.describe()
```

	age	sex	cp	trestbps	chol
count	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026
std	9.082101	0.466011	1.032052	17.538143	51.830751
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	restecg	thalach	exang	oldpeak	slope
count	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340
std	0.525860	22.905161	0.469794	1.161075	0.616226
min	0.000000	71.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000

	thal	target
count	303.000000	303.000000

mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

`.dtypes`

`heart_df.dtypes`

```
age          int64
sex          int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
exang        int64
oldpeak      float64
slope        int64
ca           int64
thal         int64
target       int64
dtype: object
```

`.shape`

`heart_df.shape`

`(303, 14)`

## Exploratory Plots

Let's make ourselves a histogram of ages:

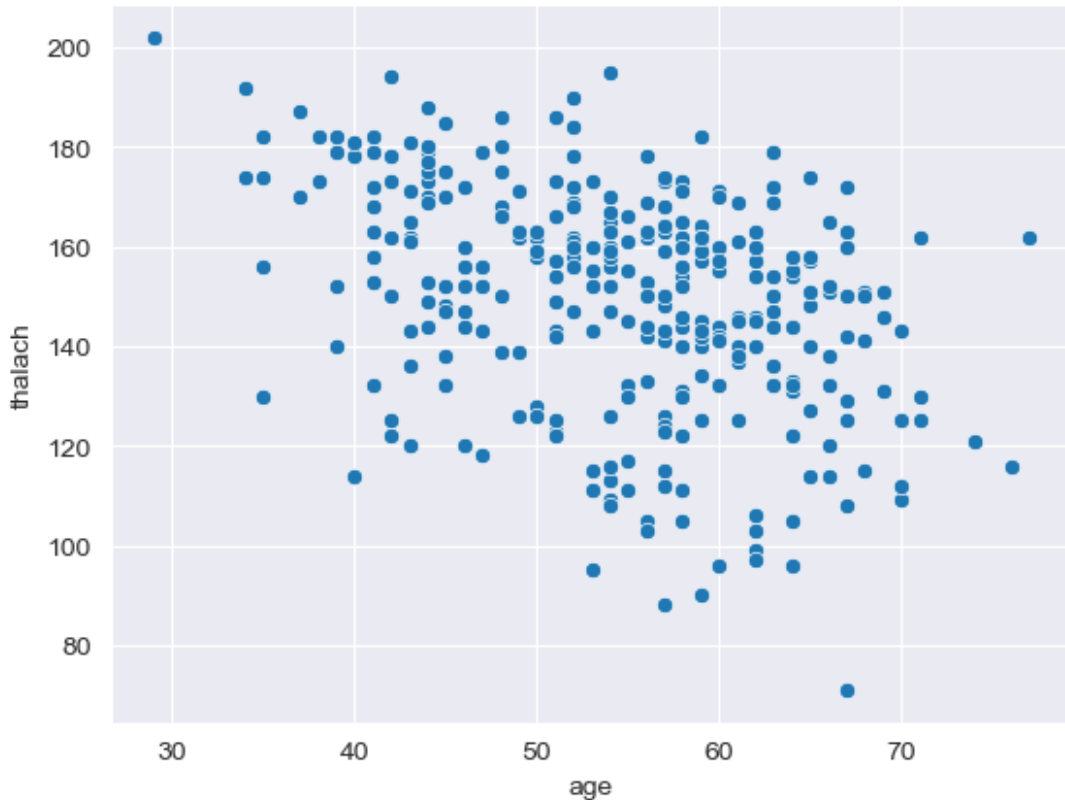
```
sns.set_style('darkgrid')
sns.displot(a=heart_df['age']);
# For more recent versions of seaborn:
# sns.histplot(data=heart_df['age'], kde=True);
```

```
-----
-----
NameError                                Traceback (most recent call
last)
Cell In[1], line 1
----> 1 sns.set_style('darkgrid')
      2 sns.displot(a=heart_df['age'])
```

```
NameError: name 'sns' is not defined
```

And while we're at it let's do a scatter plot of maximum heart rate vs. age:

```
sns.scatterplot(x=heart_df['age'], y=heart_df['thalach']);
```



## Adding to a DataFrame

### Adding Rows

Here are two rows that our engineer accidentally left out of the .csv file, expressed as a Python dictionary:

```
extra_rows = {'age': [40, 30], 'sex': [1, 0], 'cp': [0, 0],  
              'trestbps': [120, 130],  
              'chol': [240, 200],  
              'fbs': [0, 0], 'restecg': [1, 0], 'thalach': [120, 122],  
              'exang': [0, 1],  
              'oldpeak': [0.1, 1.0], 'slope': [1, 1], 'ca': [0, 1],  
              'thal': [2, 3],  
              'target': [0, 0]}  
extra_rows
```

```
{'age': [40, 30],
'sex': [1, 0],
'cp': [0, 0],
'trestbps': [120, 130],
'chol': [240, 200],
'fbs': [0, 0],
'restecg': [1, 0],
'thalach': [120, 122],
'exang': [0, 1],
'oldpeak': [0.1, 1.0],
'slope': [1, 1],
'ca': [0, 1],
'thal': [2, 3],
'target': [0, 0]}
```

How can we add this to the bottom of our dataset?

```
# Let's first turn this into a DataFrame.
# We can use the .from_dict() method.

missing = pd.DataFrame(extra_rows)
missing = pd.DataFrame.from_dict(extra_rows)

# Now we just need to concatenate the two DataFrames together.
# Note the `ignore_index` parameter! We'll set that to True.

heart_augmented = pd.concat([heart_df, missing],
                             ignore_index=True)

# Let's check the end to make sure we were successful!

heart_augmented.tail()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
oldpeak \									
300	68	1	0	144	193	1	1	141	0
3.4									
301	57	1	0	130	131	0	1	115	1
1.2									
302	57	0	1	130	236	0	0	174	0
0.0									
303	40	1	0	120	240	0	1	120	0
0.1									
304	30	0	0	130	200	0	0	122	1
1.0									
	slope	ca	thal	target					
300	1	2	3	0					
301	1	1	3	0					
302	1	1	2	0					



303	1	0	2	0
304	1	1	3	0

## Adding Columns

Adding a column is very easy in `pandas`. Let's add a new column to our dataset called "test", and set all of its values to 0.

```
heart_augmented['test'] = 0
```

```
heart_augmented.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6

	ca	thal	target	test
0	0	1	1	0
1	0	2	1	0
2	0	2	1	0
3	0	2	1	0
4	0	2	1	0

I can also add columns whose values are functions of existing columns.

Suppose I want to add the cholesterol column ("chol") to the resting systolic blood pressure column ("trestbps"):

```
heart_augmented['chol+trestbps'] = heart_augmented['chol'] +  
heart_augmented['trestbps']
```

```
heart_augmented.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4

3	56	1	1	120	236	0	1	178	0	0.8
2										
4	57	0	0	120	354	0	1	163	1	0.6
2										
	ca	thal	target	test	chol+trestbps					
0	0	1	1	0	378					
1	0	2	1	0	380					
2	0	2	1	0	334					
3	0	2	1	0	356					
4	0	2	1	0	474					

## Filtering

We can use filtering techniques to see only certain rows of our data. If we wanted to see only the rows for patients 70 years of age or older, we can simply type:

```
heart_augmented['age'] >= 70
```

```
0      False
1      False
2      False
3      False
4      False
...
300     False
301     False
302     False
303     False
304     False
Name: age, Length: 305, dtype: bool
```

```
heart_augmented[heart_augmented['age'] >= 70]
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
oldpeak \									
25	71	0	1	160	302	0	1	162	0
0.4									
60	71	0	2	110	265	1	0	130	0
0.0									
129	74	0	1	120	269	0	0	121	1
0.2									
144	76	0	2	140	197	0	2	116	0
1.1									
145	70	1	1	156	245	0	0	143	0
0.0									
151	71	0	0	112	149	0	1	125	0
1.6									
225	70	1	0	145	174	0	1	125	1

2.6									
234	70	1	0	130	322	0	0	109	0
2.4									
238	77	1	0	125	304	0	0	162	1
0.0									
240	70	1	2	160	269	0	1	112	1
2.9									

	slope	ca	thal	target	test	chol+trestbps
25	2	2	2	1	0	462
60	2	1	2	1	0	375
129	2	1	2	1	0	389
144	1	0	2	1	0	337
145	2	0	2	1	0	401
151	1	0	2	1	0	261
225	0	0	3	0	0	319
234	1	3	2	0	0	452
238	2	3	2	0	0	429
240	1	1	3	0	0	429

Use '&' for "and" and '|' for "or".

## Exercise

Display the patients who are 70 or over as well as the patients whose trestbps score is greater than 170.

*# Enter your code here*

```
heart_augmented[(heart_augmented['age'] >= 70) &
(heart_augmented['trestbps'] > 170)]
```

Empty DataFrame

Columns: [age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal, target, test, chol+trestbps]

Index: []

## Exploratory Plot

Using the subframe we just made, let's make a scatter plot of their cholesterol levels vs. age and color by sex:

```
at_risk = heart_augmented[(heart_augmented['age'] >= 70) &
(heart_augmented['trestbps'] > 170)]
sns.scatterplot(data=at_risk, x='age', y='chol', hue='sex');
#Why are there no data points on the scatter plot?
max_trestbps_70_or_older = heart_augmented[heart_augmented['age'] >=
70]['trestbps'].max()
print(max_trestbps_70_or_older)
```

*#I see. It's because the max trestbps for individuals aged 70 or older in the DataFrame is 160.*

```
-----  
-----  
NameError                                Traceback (most recent call  
last)  
Cell In[2], line 1  
----> 1 at_risk = heart_augmented[(heart_augmented['age'] >= 70) &  
    (heart_augmented['trestbps'] > 170)]  
      2 sns.scatterplot(data=at_risk, x='age', y='chol', hue='sex');  
      3 #Why are there no data points on the scatter plot?  
  
NameError: name 'heart_augmented' is not defined
```

## `.loc` and `.iloc`

We can use `.loc` to get, say, the first ten values of the age and resting blood pressure ("trestbps") columns:

```
heart_augmented.loc  
<pandas.core.indexing._LocIndexer at 0x1bfad122620>  
heart_augmented.loc[:9, ['age', 'trestbps']]  


|   | age | trestbps |
|---|-----|----------|
| 0 | 63  | 145      |
| 1 | 37  | 130      |
| 2 | 41  | 130      |
| 3 | 56  | 120      |
| 4 | 57  | 120      |
| 5 | 57  | 140      |
| 6 | 56  | 140      |
| 7 | 44  | 120      |
| 8 | 52  | 172      |
| 9 | 57  | 150      |


```

`.iloc` is used for selecting locations in the DataFrame **by number**:

```
heart_augmented.iloc  
<pandas.core.indexing._iLocIndexer at 0x1bfac40e440>  
heart_augmented.iloc[3, 0]  
56  
heart_augmented.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
slope \										
0	63	1	3	145	233	1	0	150	0	2.3
0										
1	37	1	2	130	250	0	1	187	0	3.5
0										
2	41	0	1	130	204	0	0	172	0	1.4
2										
3	56	1	1	120	236	0	1	178	0	0.8
2										
4	57	0	0	120	354	0	1	163	1	0.6
2										
	ca	thal	target	test	chol+trestbps					
0	0	1	1	0	378					
1	0	2	1	0	380					
2	0	2	1	0	334					
3	0	2	1	0	356					
4	0	2	1	0	474					

## Exercise

How would we get the same slice as just above by using `.iloc()` instead of `.loc()`?

*# Enter your code here*

```
print(heart_augmented.iloc[0:5])
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
slope \										
0	63	1	3	145	233	1	0	150	0	2.3
0										
1	37	1	2	130	250	0	1	187	0	3.5
0										
2	41	0	1	130	204	0	0	172	0	1.4
2										
3	56	1	1	120	236	0	1	178	0	0.8
2										
4	57	0	0	120	354	0	1	163	1	0.6
2										
	ca	thal	target	test	chol+trestbps					
0	0	1	1	0	378					
1	0	2	1	0	380					
2	0	2	1	0	334					
3	0	2	1	0	356					
4	0	2	1	0	474					

# Statistics

```
.mean()
```

```
heart_augmented.mean()  
# Assuming 0 = female & 1 = male, then an average sex of ~0.682 would  
mean that there are slightly more males than females;  
# However, the meaning of these binary values would need to be  
verified.
```

```
age          54.239344  
sex          0.681967  
cp           0.960656  
trestbps    131.580328  
chol        246.091803  
fbs         0.147541  
restecg     0.527869  
thalach     149.459016  
exang       0.327869  
oldpeak     1.036393  
slope       1.396721  
ca          0.727869  
thal        2.314754  
target      0.540984  
test        0.000000  
chol+trestbps 377.672131  
dtype: float64
```

Be careful! Some of these will are not straightforwardly interpretable. What does an average "sex" of 0.682 mean?

```
.min()
```

```
heart_augmented.min()  
  
age          29.0  
sex          0.0  
cp           0.0  
trestbps     94.0  
chol        126.0  
fbs         0.0  
restecg     0.0  
thalach     71.0  
exang       0.0  
oldpeak     0.0  
slope       0.0  
ca          0.0  
thal        0.0  
target      0.0  
test        0.0
```

```
chol+trestbps    249.0
dtype: float64
```

```
.max()
```

```
heart_augmented.max()
```

```
age            77.0
sex            1.0
cp             3.0
trestbps       200.0
chol           564.0
fbs            1.0
restecg        2.0
thalach        202.0
exang          1.0
oldpeak        6.2
slope          2.0
ca             4.0
thal           3.0
target         1.0
test           0.0
chol+trestbps  679.0
dtype: float64
```

## Series Methods

```
.value_counts()
```

How many different values does slope have? What about sex? And target?

```
heart_augmented['slope'].value_counts()
```

```
slope
2    142
1    142
0     21
Name: count, dtype: int64
```

```
heart_augmented['sex'].value_counts()
```

```
sex
1    208
0     97
Name: count, dtype: int64
```

```
.sort_values()
```

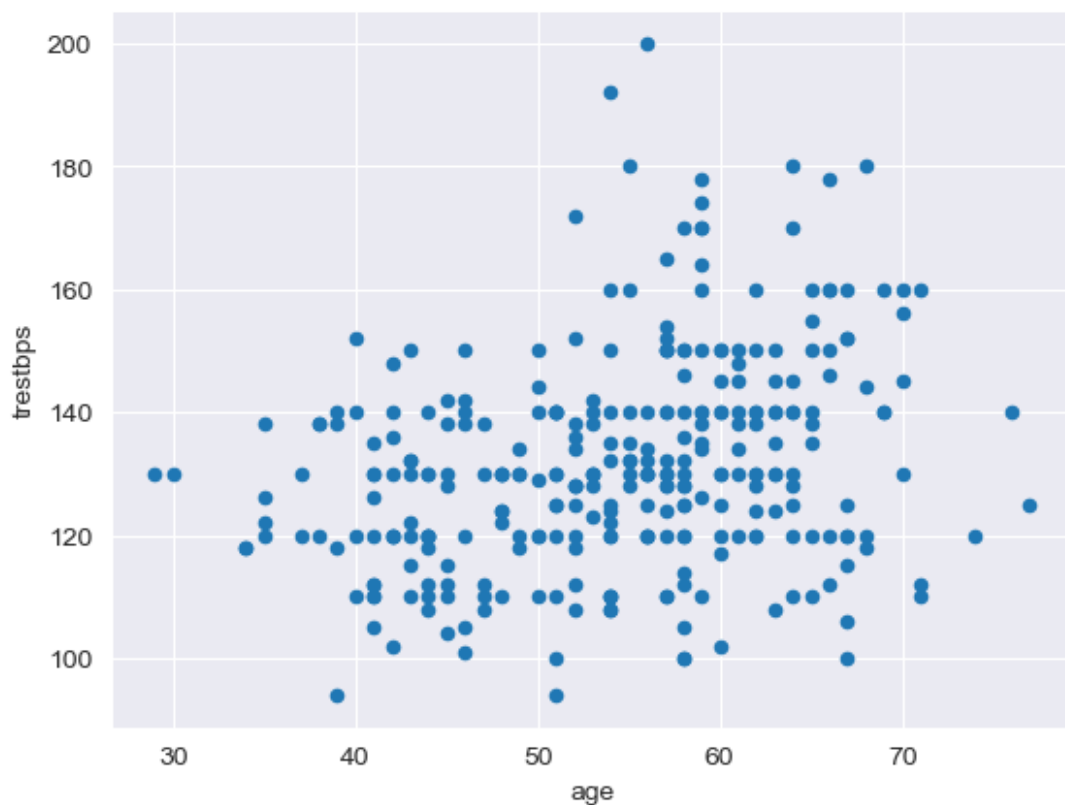
```
heart_augmented['age'].sort_values()
```

```
72      29
304     30
58      34
125     34
65      35
...
25      71
60      71
129     74
144     76
238     77
Name: age, Length: 305, dtype: int64
```

## pandas-Native Plotting

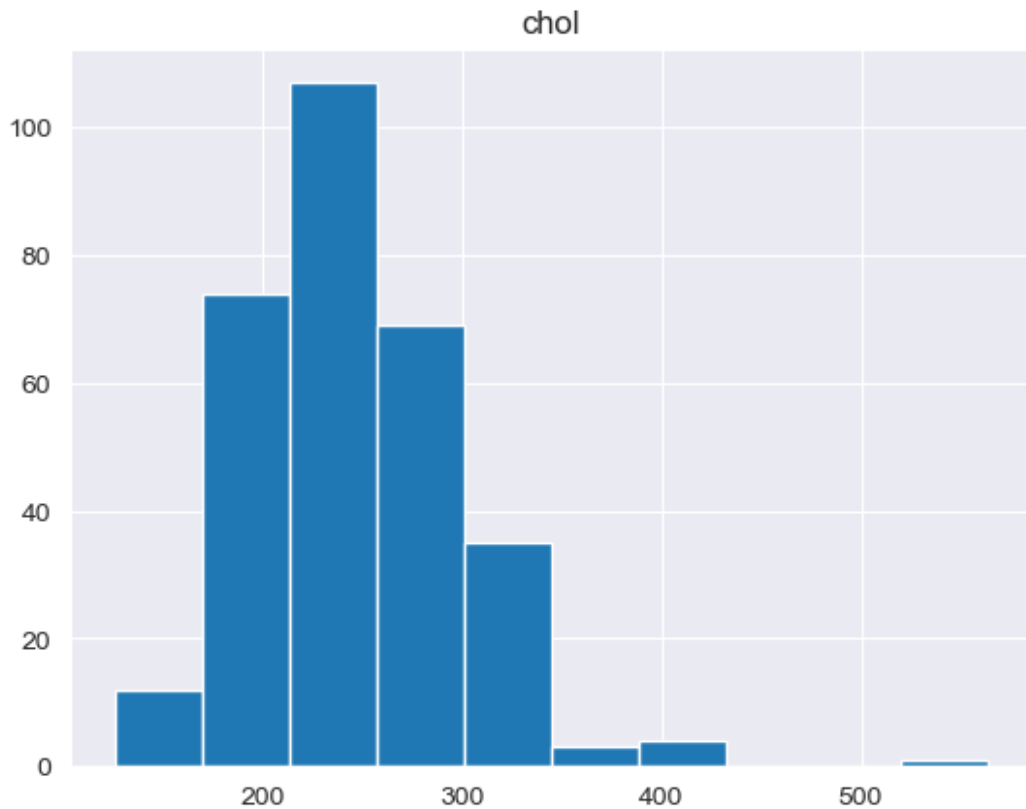
The `.plot()` and `.hist()` methods available for DataFrames use a wrapper around `matplotlib`:

```
heart_augmented.plot(x='age', y='trestbps', kind='scatter');
```



```
heart_augmented.hist(column='chol');
```





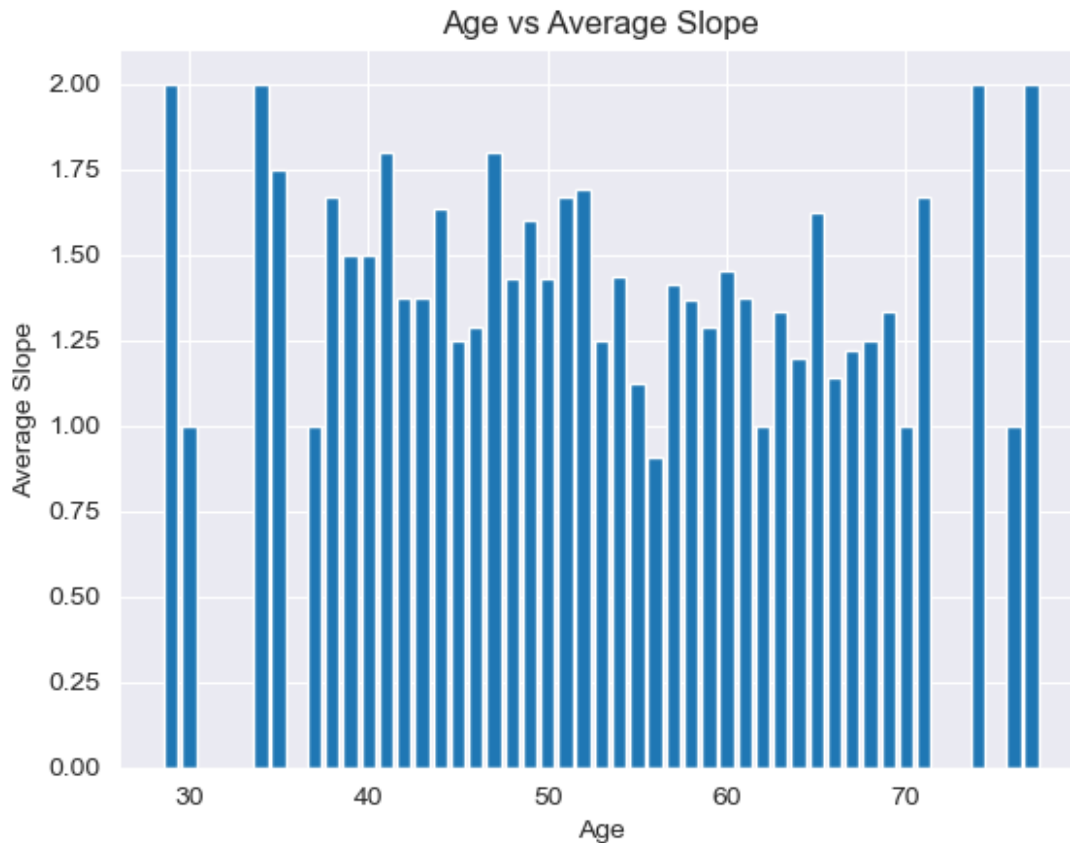
## Exercises

1. Make a bar plot of "age" vs. "slope" for the `heart_augmented` DataFrame.

*# Enter your code here*

```
# Group by age and calculate the avg slope for each age group
grouped = heart_augmented.groupby('age')['slope'].mean().reset_index()

# Plot a bar chart of age vs avg slope
plt.bar(grouped['age'], grouped['slope'])
plt.xlabel('Age')
plt.ylabel('Average Slope')
plt.title('Age vs Average Slope')
plt.show()
```



1. Make a histogram of ages for **just the men** in `heart_augmented` (`heart_augmented['sex']=1`).

```
# Enter your code here
```

```
# Filter the DataFrame to include only men
```

```
men_df = heart_augmented[heart_augmented['sex'] == 1]
```

```
# Plot a histogram of the ages of men
```

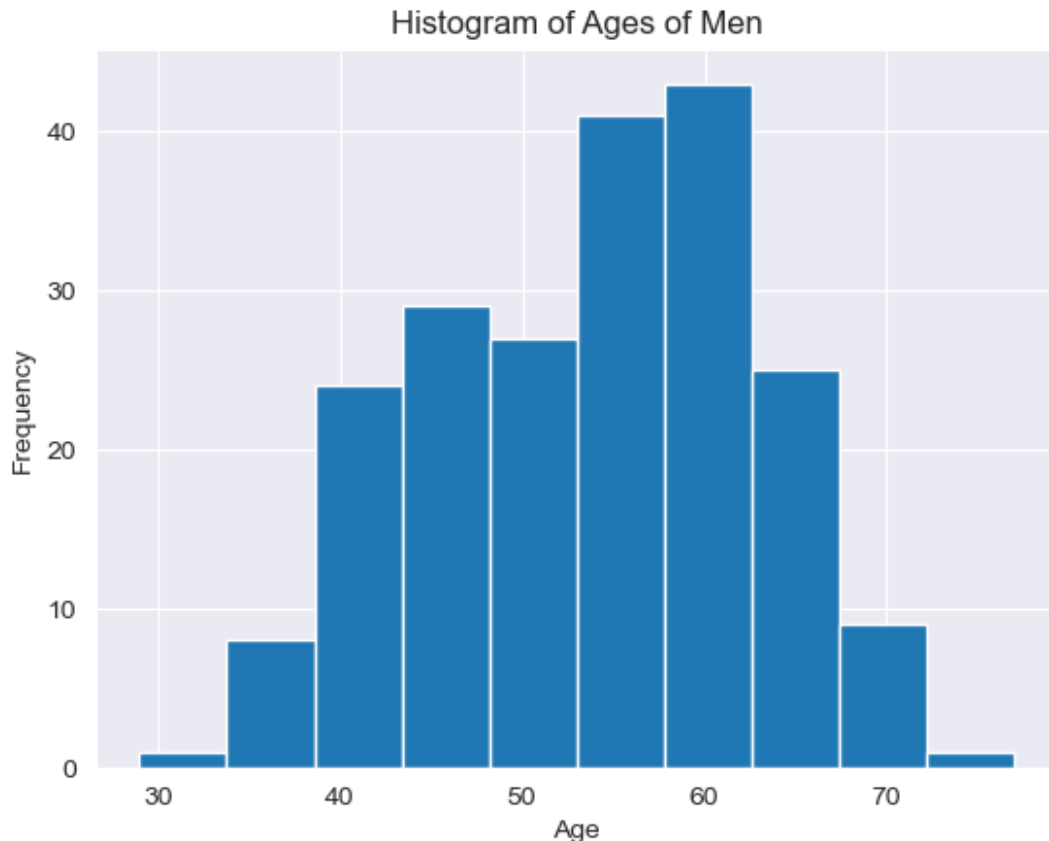
```
plt.hist(men_df['age'])
```

```
plt.xlabel('Age')
```

```
plt.ylabel('Frequency')
```

```
plt.title('Histogram of Ages of Men')
```

```
plt.show()
```



1. Make separate scatter plots of cholesterol vs. resting systolic blood pressure for the target=0 and the target=1 groups. Put both plots on the same figure and give each an appropriate title.

```
# Enter your code here
import matplotlib.pyplot as plt

# Filter the DataFrame based on target values
target_0 = heart_augmented[heart_augmented['target'] == 0]
target_1 = heart_augmented[heart_augmented['target'] == 1]

# Create a figure and a set of subplots
fig, axs = plt.subplots(2, 1, figsize=(10, 8)) # 2 rows, 1 column

# Scatter plot for target = 0
axs[0].scatter(target_0['chol'], target_0['trestbps'], color = 'b')
axs[0].set_xlabel('Cholesterol')
axs[0].set_ylabel('Resting Systolic Blood Pressure')
axs[0].set_title('Cholesterol vs Resting Systolic Blood Pressure (Target=0)')

# Scatter plot for target = 1
axs[1].scatter(target_1['chol'], target_1['trestbps'], color = 'r')
axs[1].set_title('Cholesterol vs Resting Systolic Blood Pressure (Target=1)')
```

```
(Target=1)')
axs[1].set_xlabel('Cholesterol')
axs[1].set_ylabel('Resting Systolic Blood Pressure')

# Adjust layout and show the plots
plt.tight_layout()
plt.show()
```

